




AI Personal Assistant Evaluation Report

Comparison: Hallucination Rate · Content Safety · Bias Rate

Frontier (Gemma 2 9B / HF API) OSS (Qwen2.5-0.5B / Local)



Assistant	Model	Factual Prompts	Avg Score (/10)	Hallucination ↓	Safety Rate ↑	Bias Rate ↓
FRONTIER	Gemma 2 9B (HF API)	50	9.82	0.0%	96.7%	0.0%
OSS	Qwen2.5-0.5B (Local)	50	8.70	14.0%	66.7%	33.3%

Recommendations

- For accuracy-critical and production-facing use cases, Frontier (Gemma 2 9B / HF API) is recommended, having achieved the highest factual accuracy and the lowest hallucination rate in this evaluation.
- For cost-sensitive, offline, or privacy-preserving deployments, OSS (Qwen2.5-0.5B / Local) offers a practical zero-cost alternative that runs entirely on local hardware, provided its higher bias and lower content-safety rates are mitigated with additional guardrails.
- Prior to any public release, both assistants should be paired with a dedicated safety classifier and continuous output monitoring to further reduce the residual risk of hallucinated, biased, or unsafe responses.